

Final Project new version

September 20, 2024

1 STAT 201 Project

1.1 Group 11

Ray Tao
Florence Wang
Lesley Mai
Harsha Chhetri

2 Compare the delay rate of two popular airlines in the United States

2.1 1. Introduction

2.1.1 1.1 Basic Information

In the past several decades, air travel has become one of the most popular options due to the rising demand for rapid travel. The most depressing aspect of a flight delay is that it is occasionally unavoidable. Recent research has shown that among different airlines, about 0% to 34% were delayed (FlightAware, 2023), and the average delay rate of an airline has become an essential standard for customers to evaluate and choose a company. Therefore, it is crucial to investigate and compare the delay rates of major airlines to understand how they differ and identify potential areas for improvement. We will specifically look into any differences in delay rates between two industry giants, American Airlines and Delta Airlines as well as potential causes for these variations. Our study seeks to provide valuable insights for both airlines and customers and contribute to the ongoing efforts to improve the air travel experience.

2.1.2 1.2 Research Question

Is there a significant difference between the delay rate of flights operated by Delta Airlines and American Airlines?

To answer this question, we have chosen a publicly accessible dataset created by Ulrik Thyge Pedersen (2023), which includes data from 7.81k flights from five different airline companies on the measures that we are interested in: DL & AA airlines, and delay or not. We will be answering our question by first looking at the proportion of delayed flights of these two airlines and how they differ based on the categorical variables.

2.2 2. Methods and Results

2.2.1 2.1 Preliminary Data Analysis

```
[1]: # Needed packages
library(tidyverse)
library(infer)
library(dplyr)
library(readr)
library(broom)
library(ggplot2)
library(cowplot)
set.seed(201)
```

```
Attaching packages: tidyverse
1.3.2
ggplot2 3.3.6 purrr 0.3.4
tibble 3.1.8 dplyr 1.0.10
tidyr 1.2.1 stringr 1.4.1
readr 2.1.2 forcats 0.5.2

Conflicts:
tidyverse_conflicts()
dplyr::filter() masks stats::filter()
dplyr::lag() masks stats::lag()
```

Load Data Our dataset is formatted as a .csv file with headers. We will use `read_csv()` to load our raw data in for wrangling.

```
[2]: airlines_delay <- read_csv("https://raw.githubusercontent.com/rayyyy122/
STAT-201-Project/main/airlines_delay.csv")
```

```
Rows: 539382 Columns: 8
Column specification
```

```
Delimiter: ","
chr (3): Airline, AirportFrom, AirportTo
dbl (5): Flight, Time, Length, DayOfWeek, Class
```

Use ``spec()`` to retrieve the full column specification for this data.

Specify the column types or set ``show_col_types = FALSE`` to quiet this message.

Raw Dataset To help gauge what wrangling is required, let's use `head()`, `tail()` to preview the first and last 3 rows of our data.

```
[ ]: airlines_delay |> head(3)
airlines_delay |> tail(3)
```

Table 1: Raw Airlines Delay Dataset

Clean and Wrangle Data To clean and wrangle our data, we are going to:

1. Select the relevant columns for our question
2. Filter out rows containing NA

Since our original data has more than 100000 rows of observations, which makes bootstrapping operations run very slowly, we use `rep_sample_n` to draw a random sample from our data with a size of 2000.

```
[12]: # clean and wrangle data
delay <- airlines_delay |>
  filter(Airline %in% c("DL", "AA")) |>
  rename(Delay = Class) |>
  filter(!is.na(Airline)) |>
  filter(!is.na(Delay)) |>
  mutate(Delay = as.factor(Delay)) |> # 0 means not Delayed
  select(c(Airline, Delay))

# randomly draw a sample with size of 2000
delay_data <- delay |>
  rep_sample_n(reps = 1, size = 2000)

head(delay_data)
```

	replicate	Airline	Delay
	<int>	<chr>	<fct>
A grouped_df: 6 × 3	1	AA	1
	1	DL	0
	1	DL	0
	1	DL	1
	1	DL	0
	1	DL	0

Table 2: DL Airline and AA Airline Delay or Not

Broad Overview of Data In this section, we will compute the estimates of the parameter “proportion” of delayed flights for each airline first, and plot the relevant data. We used different colors for delayed flights and proper flights to observe the difference more clearly.

```
[5]: # Compute the estimates of the parameter
delay_data_summary <- delay_data %>%
  group_by(Airline, Delay) %>%
  summarize(count = n()) %>%
  group_by(Airline) %>%
  mutate(total = sum(count),
         proportion = count/total)
```

```

delay_data_summary

# calculate observed test statistic
obs_diff_prop <- delay_data |>
  specify(Delay ~ Airline, success = "1") |>
  calculate(stat = "diff in props", order = c("DL", "AA")) |>
  rename(observed_test_stat = stat)

obs_diff_prop

```

`summarise()` has grouped output by 'Airline'. You can override using the
`.groups` argument.

	Airline	Delay	count	total	proportion
	<chr>	<fct>	<int>	<int>	<dbl>
A grouped_df: 4 × 5	AA	0	534	849	0.6289753
	AA	1	315	849	0.3710247
	DL	0	615	1151	0.5343180
	DL	1	536	1151	0.4656820

	observed_test_stat
	<dbl>
A infer: 1 × 1	0.09465728

We create two visualizations of the data: one displays the count of delayed versus not delayed flights of each airline, and another shows the proportion of delayed versus not delayed flights of each airline.

```

[13]: plot <- delay_data %>%
  ggplot(aes(x = Airline, fill = Delay)) +
  geom_bar(width = 0.5) +
  ggtitle("Count of Delayed Flights by Airline") +
  labs(y = "Number of Flights", caption = "Figure 1. Visualize the number of
  ↪ delayed flights of two airlines") +
  theme(plot.caption = element_text(size = 13, color = "black", hjust = 0,
  ↪ face = "italic")) +
  scale_fill_discrete(name = "Delay or Not",
    breaks = c("0", "1"),
    labels = c("No", "Yes")) +
  annotate("text", x = 1, y = 20, label = "Delayed Rate: 37.1%", col =
  ↪ "purple", size = 4) +
  annotate("text", x = 2, y = 20, label = "Delayed Rate: 46.5%", col =
  ↪ "purple", size = 4)
plot

```

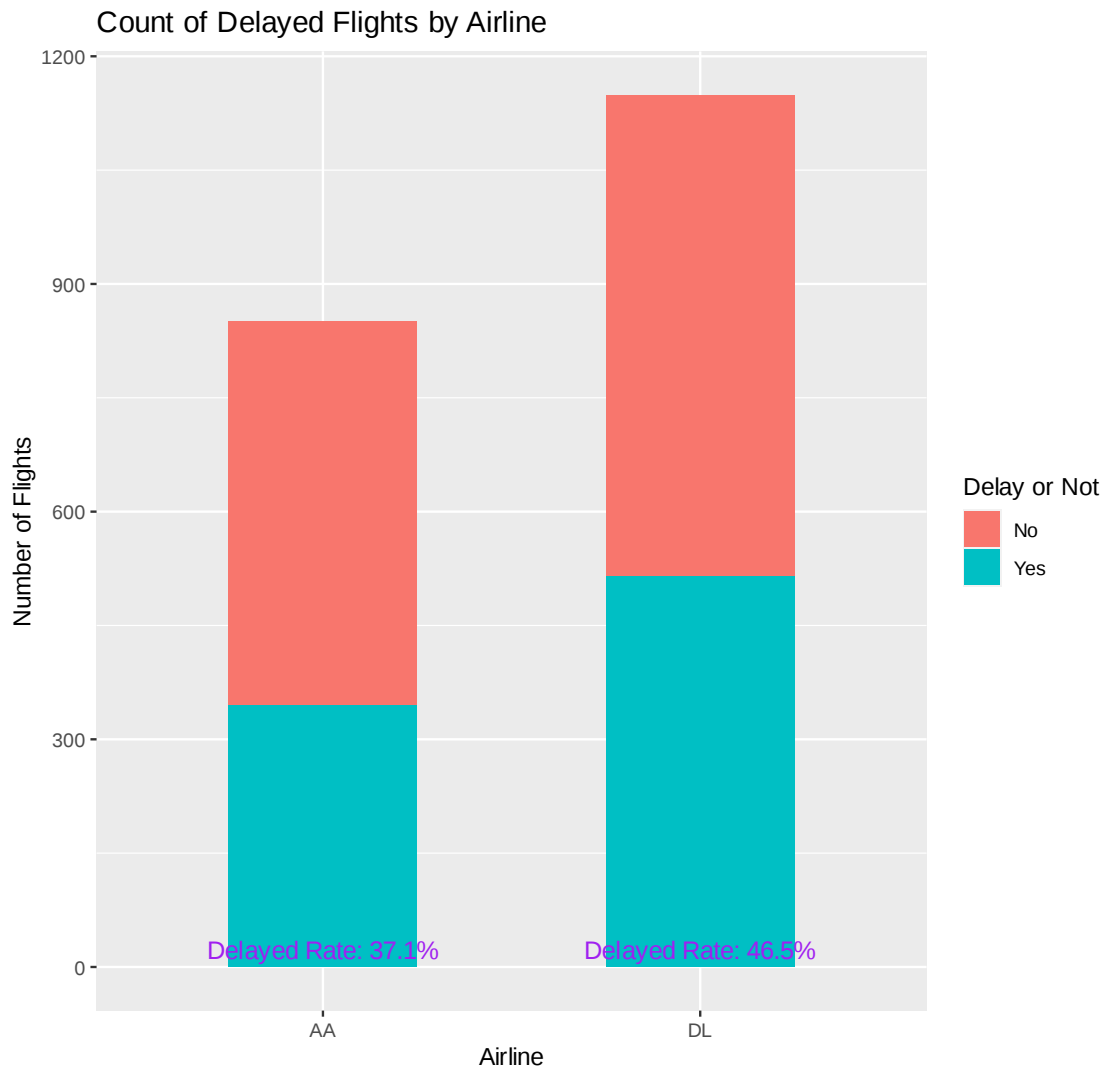


Figure 1. Visualize the number of delayed flights of two airlines

```
[15]: plot2 <- delay_data %>%
  ggplot(aes(x = Airline, fill = Delay)) +
  geom_bar(width = 0.5, position = "fill") +
  ggtitle("Proportion of Delayed Flights by Airline") +
  labs(y = "Number of Flights", caption = "Figure 2. Visualize the proportion
of delayed flights of two airlines") +
  theme(plot.caption = element_text(size = 13, color = "black", hjust = 0,
face = "italic")) +
  scale_fill_discrete(name = "Delay or Not",
                      breaks = c("0", "1"),
                      labels = c("No", "Yes")) +
  annotate("text", x = 1, y = 0.02, label = "Delayed Rate: 37.1%", col =
"purple", size = 4) +
```

```

annotate("text", x = 2, y = 0.02, label = "Delayed Rate: 46.5%", col = "purple", size = 4)

```

plot2

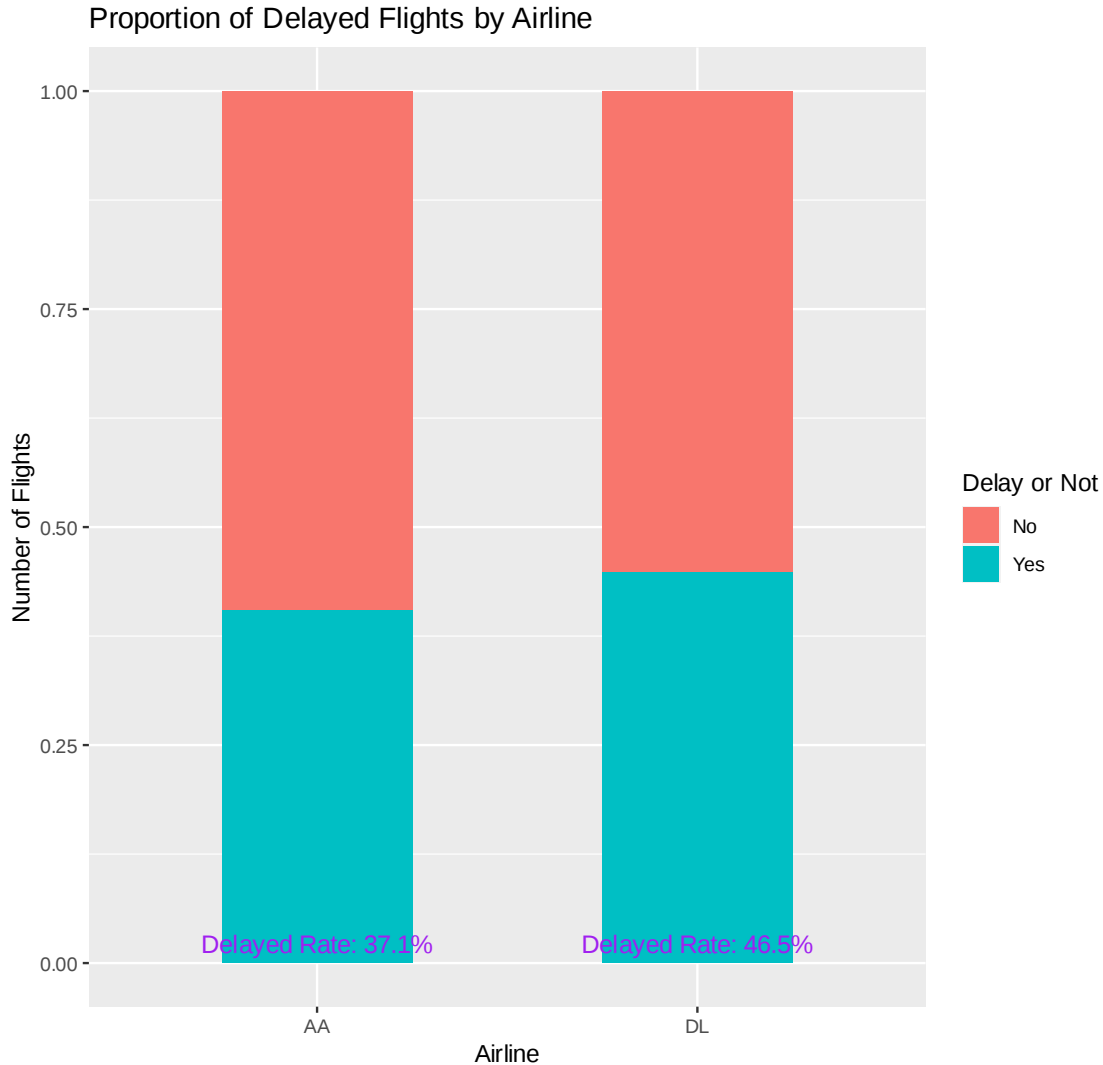


Figure 2. Visualize the proportion of delayed flights of two airlines

2.2.2 2.2 Methods overview

This project's primary objective is to investigate and compare the proportion of delayed flights operated by Delta Airlines and American Airlines, and to test the hypothesis of whether there is a statistically significant difference between the two proportions. To achieve this, we will employ two different statistical methods: **bootstrapping** and **theory-based approaches**.

In the bootstrapping method, we will use the **infer** package to perform a permutation test. Specifically, we will randomly resample the data with replacement to create a null model, where there

is no difference between the proportions of delayed flights for Delta and American Airlines. We will then calculate the difference in proportions between the two airlines and compare it to the observed difference. If the observed difference is outside the range of the null distribution, we can reject the null hypothesis and conclude that there is a statistically significant difference between the two proportions.

For the theory-based approach, we will use the Central Limit Theorem two-sample z-test. This method assumes that the proportions of delayed flights for both airlines follow a normal distribution and the sample size is large enough to apply the Central Limit Theorem. We will calculate the test statistic, which measures the difference between the two proportions, and compare it to the p-value of the z-distribution at a significance level of $\alpha = 0.05$. If the test statistic is larger than the critical value, we can reject the null hypothesis and conclude that there is a statistically significant difference between the two proportions.

Our null and alternative hypothesis is:

- $H_0 : p_1 = p_2$
- $H_a : p_1 \neq p_2$

where:

p_1 is the proportion of delayed flights in all flights operated by Delta Airlines

p_2 is the proportion of delayed flights in all flights operated by American Airlines

Based on the sample proportions and the graph above, we expect to reject the null hypothesis and accept the alternative that the delay rate of these two airlines is different.

2.2.3 2.3 Hypothesis test via bootstrapping method

We applied bootstrapping asymptotic method in the `infer` package to create the null model.

```
[8]: # create the null model
null_distribution <- delay_data |>
  specify(formula = Delay ~ Airline, success = "1") |>
  hypothesize(null = "independence") |>
  generate(reps = 500, type = "permute") |>
  calculate(stat = "diff in props", order = c("DL", "AA"))

[9]: # visualize the null model and the position of observed test statistic
visualize(null_distribution, bins = 10) +
  shade_p_value(obs_stat = obs_diff_prop, direction = "both") +
  labs(x = "Difference in Proportion", caption = "Figure 3. Visualize the null
distribution and the position of observed test statistic") +
  theme(plot.caption = element_text(size = 13, color = "black", hjust = 0,
face = "italic"))

# calculate the P-value
get_p_value(null_distribution, obs_stat = obs_diff_prop, direction = "both")
```

Warning message:

"Please be cautious in reporting a p-value of 0. This result is an approximation based on the number of `reps` chosen in the `generate()` step. See `?get_p_value()` for more information."

A tibble: 1 × 1

p_value
0




Figure 3. Visualize the null distribution and the position of observed test statistic

We then created the bootstrapping model, calculated the 95% confidence interval and visualized the model shaded with the range of the confidence interval.

```
[10]: # calculate the 95% confidence interval
bootstrap_distribution <- delay_data |>
```



```

specify(formula = Delay ~ Airline, success = "1") |>
hypothesize(null = "independence") |>
generate(reps = 500, type = "bootstrap") |>
calculate(stat = "diff in props", order = c("DL", "AA"))

ci <- bootstrap_distribution |>
  get_confidence_interval(type = "se", point_estimate = obs_diff_prop, level_
↪= 0.95)
ci

# visualize the confidence interval
ci_plot <- visualize(bootstrap_distribution, bins = 10) +
  shade_confidence_interval(endpoints = ci) +
  labs(x = "Difference in Proportion", caption = "Figure 4. Visualize the_
↪confidence interval") +
  annotate("text", x = 0.15, y = 100, label = "Upper Tail: \n 0.1362") +
  annotate("text", x = 0.04, y = 100, label = "Lower Tail: \n 0.0531") +
  theme(plot.caption = element_text(size = 13, color = "black", hjust = 0,
↪face = "italic"))
ci_plot

```

	lower_ci	upper_ci
A tibble: 1 × 2	<dbl>	<dbl>
	0.05309204	0.1362225

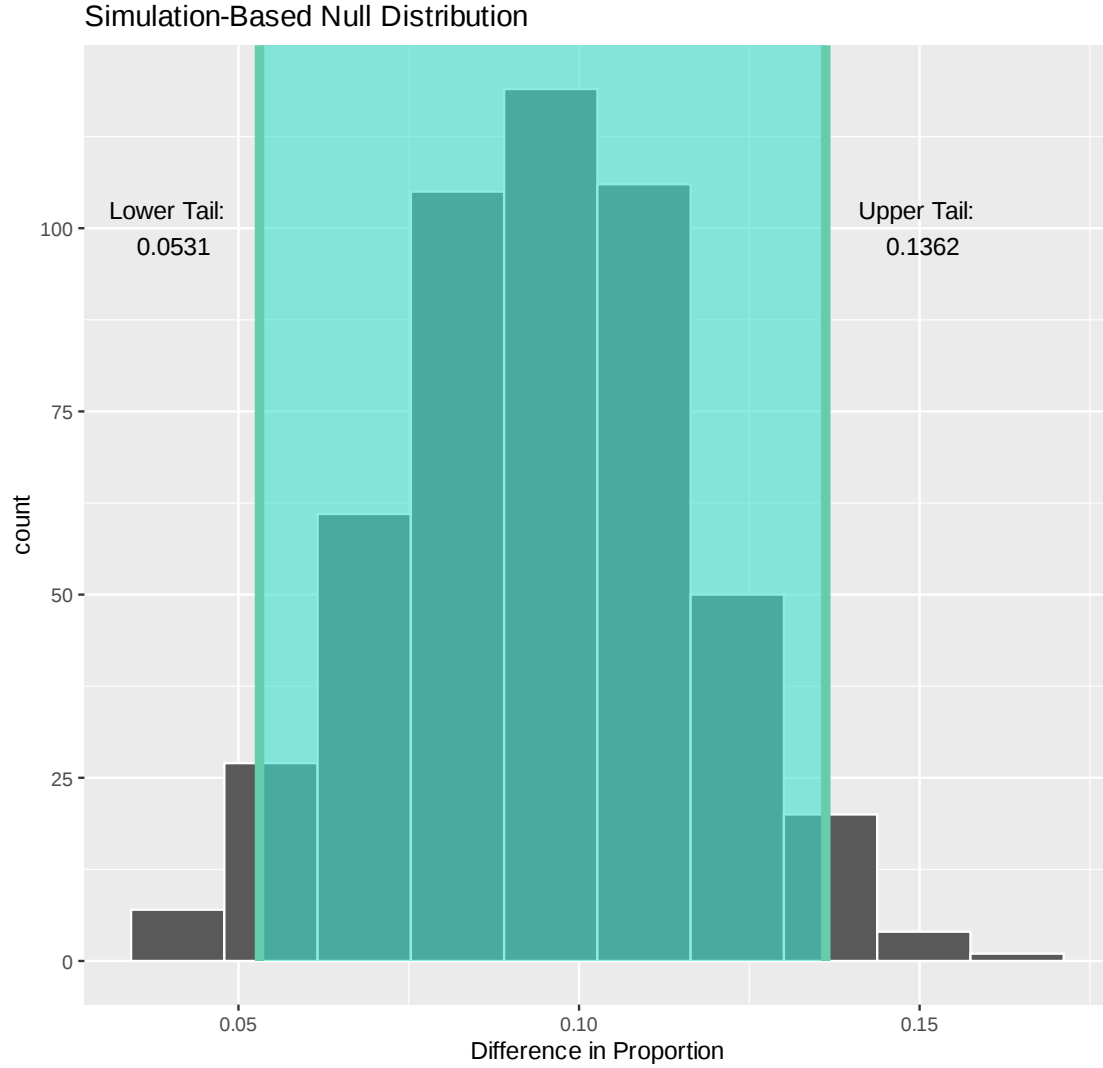


Figure 4. Visualize the confidence interval

Result First, we notice that the p-value is equal to 0 and the position of the test statistic covers none of the observations in the null distribution. Based on these two discoveries, we decided to reject the null hypothesis at a common 95% significance level. In other words, we believe that there exists a difference between the proportion of delayed flights operated by Delta Airlines versus American Airlines.

2.2.4 2.4 Hypothesis test via theory based approach

- $H_0 : p_1 = p_2$
- $H_a : p_1 \neq p_2$

$$\text{test statistic} = \frac{\hat{p}_1 - \hat{p}_2}{\sqrt{\hat{p}(1-\hat{p})\left(\frac{1}{n_1} + \frac{1}{n_2}\right)}}$$

where $\hat{p} = \frac{n_1\hat{p}_1 + n_2\hat{p}_2}{n_1 + n_2}$

\hat{p}_1 : the proportion of delayed flights operated by Delta Airlines

\hat{p}_2 : the proportion of delayed flights operated by American Airlines

n_1 : the number of flights operated by Delta Airlines in the sample

n_2 : the number of flights operated by American Airlines in the sample

```
[16]: n1 <- sum(delay_data$Airline == "DL")
n2 <- sum(delay_data$Airline == "AA")
phat1 <- sum(delay_data$Airline == "DL" & delay_data$Delay == "1") / n1
phat2 <- sum(delay_data$Airline == "AA" & delay_data$Delay == "1") / n2

# perform the hypothesis test
result <- tidy(
  prop.test(
    x = c(n1 * phat1, n2 * phat2),
    n = c(n1, n2),
    correct = FALSE))

result
```

	estimate1	estimate2	statistic	p.value	parameter	conf.low	conf.high	method
A tibble: 1 × 9	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<chr>
	0.4482158	0.4042303	3.860293	0.04944169	1	0.0002365549	0.08773449	2-sam

Result From the tidy output of the theory-based hypothesis test, we can see that the P value is 2.317e-05 which is much less than the significance level of 0.05. So we can reject the null hypothesis that there is no difference between the proportion of delayed flights between Delta Airlines and American Airlines. To be more specific, we'd like to accept the alternative hypothesis that there exists an obvious difference between them.

2.2.5 2.5 Comparison between these two methods

Based on the two result section above, we can see that both methods give very similar results: we reject the null hypothesis that the proportions of delayed flights operated by Delta and American Airlines are identical according to the small p-value vs the common 95% significance level. Though all the assumptions of the two methods are satisfied in this study, we may prefer the theory-based approach since it gives a p-value of 2.317e-05 while the bootstrapping method outputs a p-value of 0. The P-value of 0 gives the "perfect" rejection of the null hypothesis, which seems to be not really meaningful and reasonable.

2.3 3. Discussion

In this project, our research question is Is there a difference between the delay rate of flights operated by Delta Airlines and American Airlines? We figured out this concern by performing hypothesis testing on the proportion of delayed flights operated by each of these two world-famous airlines. To make the result more plausible, we apply two types of methods. For the bootstrapping method, the sample we drew randomly from the dataset contains 2000 flights in the whole of America, which

is representative and ensures bootstrapping works well. For the theory-based one, we carefully checked that both $n * p$ and $n * (1 - p)$ are larger than 10, which means the Central Limit Theorem is satisfied. Hence the theory-based approach should also work well.

After conducting hypothesis tests via different methods with some visualizations as well as computing statistical decision threshold values, we get quite similar preliminary results. As mentioned above, we find there's a significant difference between the delay rate of flights operated by Delta Airlines and American Airlines. This is what we expected at the beginning of this study. Because in reality, the delay rate is a pretty crucial standard when evaluating an airline company. Whether a flight delay or not can reflect much information, such as the efficiency of management, the force of supervision, etc.

Convincing as this research is, we merely focused on the different delayed rates of the two airlines, without consideration of other relevant factors. Therefore, there exists a bunch of limitations regarding our study and customer may not be able to judge an airline company solely based on our study. In the future, this study can lead to more concerns. One of them can be the relationship between the difference in delayed flights in terms of certain routes. Will Delta Airlines perform better than American Airlines on certain routes or vice versa? Another limitations is that it does not tell us where this significant difference lies. We need to consider is there more test method that is more accurate to find the difference between two population other than Central Limit Theorem? Further, limited by time and resources, we mainly concentrated on these two aviation giants. Questions regarding the delayed rate of other Airlines in the U.S. or even the whole world may arise.

2.4 Reference

1. FlightAware (2023). <https://flightaware.com/live/cancelled/>
2. Pedersen, U. T. (2023). Airlines Delay. <https://www.kaggle.com/datasets/ulrikthygepedersen/airlines-delay>

[]: